

Transaction Management

Transactional Frameworks for Middleware Platform
Middleware Architecture (SCS 2314)

1. What is a Transaction?

- A transaction = change from one **stable state** to another **stable state**
- In between = **unstable intermediate states**
- Once complete → new stable state reached

```
{Initial Stable State}
|
[a(1), a(2)... a(n)]  <- sequence of actions
|
{Final Stable State}
```

Example — Bank Transfer:

```
Before:  A = 1000, B = 500   <- stable
During:  A = 500,  B = 500   <- unstable (money in air)
After:   A = 500,  B = 1000  <- stable ■
```

2. Atomic Action

- Atomic = **indivisible** — all or nothing
- Two atomic operations A and B = **mutually exclusive on dependency**
- They don't interfere with each other's outcome

Three valid execution orders:

1. Concurrently: A || B
2. A first: A --> B
3. B first: B --> A

All three -> correct outcome ■

3. Composite Action

Serialised

- A **sequence** of atomic operations one after another
- Starts with initial state → ends with final state → within finite time

```
{Initial State}
|
a(1)
|
a(2)
|
a(n)
|
{Final State}
```

Concurrent

- Multiple composite actions running **at the same time**
- They access a **shared object** → risky ■

```
Process A:  a(1) --> a(2) --> a(n)
              \
              [Shared Object]  <- both access this!
              /
```

```
Process B:  b(1) --> b(2) --> b(n)
```

- If not managed → **data corruption** ■
- This is why we need **concurrency control**

4. ACID Properties

Every transaction must follow these 4 properties:

```
A -> Atomicity      (Failure Atomicity)
C -> Consistency    (Integrity)
I -> Isolation      (Concurrency Atomicity)
D -> Durability
```

A — Atomicity

- Transaction must complete **fully** or **not at all**
- Failure midway → **rollback** to original state

```
S1 --> [Action] --> S2    <- success ■
S1 --> [Action] --> FAIL  -> rollback -> S1 ■
Final state = EITHER S1 OR S2, never in between
```

C — Consistency

- Must go from one **consistent state** to another
- Intermediate inconsistent states = **hidden inside** transaction, never visible outside

```
Consistent --> [transaction] --> Consistent
                |
                inconsistent states
                (hidden inside)
```

I — Isolation

- Multiple transactions run concurrently
- Result must equal as if they ran **one after another**

```
A || B must equal either:
-> A then B  (A ; B)
OR
-> B then A  (B ; A)
Intermediate states of one -> NOT visible to another ■
```

D — Durability

- Successful transaction = **Commit**
- After commit → changes are **permanent** even if system crashes ■

```
Transaction --> COMMIT --> permanent forever ■
Transaction --> FAIL   --> rollback, nothing saved ■
```

Managing ACID — 3 Mechanisms

ACID Property	Managed By
Atomicity	Rollback / Abort
Consistency	Rollback / Abort
Isolation	Concurrency Control (Locking)
Durability	Commit

Why Concurrency?

- ## Goal

- ## Serialisability

- | S1 (serialisable ■) | S2 (serialisable ■) | S3 (NOT serialisable ■) |
|---------------------------|---------------------------|---------------------------|
| ■■■■■■■■■■■■■■■■■■■■■■■■■ | ■■■■■■■■■■■■■■■■■■■■■■■■■ | ■■■■■■■■■■■■■■■■■■■■■■■■■ |
| T1: a = a - x | T1: a = a - x | T1: a = a - x |
| T1: b = b + x | T2: a = (1+r)a | T2: a = (1+r)a |
| T1: commit | T2: b = (1+r)b | T2: b = (1+r)b <- written |
| T2: a = (1+r)a | T2: commit | T2: commit |
| T2: b = (1+r)b | T1: b = b + x | T1: b = b + x <- LOST! ■ |
| T2: commit | T1: commit | T1: commit |

S1 -> serial order -> serialisable ■
S2 -> different order but same result as S1 -> serialisable ■
S3 -> write-write conflict on b -> T1's write lost -> NOT serialisable ■

KEY RULE: Schedule is serialisable $\leftrightarrow$ dependency graph is ACYCLIC
S3 has a cycle -> NOT serialisable ■

6. Locking

- Mechanism to **delay** a transaction's progress
- Enforces serialisability
- Most frequently used technique

Two Types of Locks

Lock Type	Used For	Who can hold it
Shared	Read only	Multiple transactions ■
Exclusive	Read + Write	ONE transaction only ■

Lock Commands

```
lock_exclusive(o)  -> acquire exclusive lock on object o
lock_shared(o)     -> acquire shared lock on object o
unlock(o)          -> release lock on object o
```

Rules

- Acquire lock **before** using shared object
- Shared lock → read only | Exclusive lock → read + write
- Cannot lock already locked object (except upgrade shared → exclusive)
- Must release **all locks** after commit or abort

```
BEFORE access --> acquire lock
DURING access --> lock held
AFTER commit  --> release lock
```

7. Two Phase Locking (2PL)

2PL Rule: After releasing a lock → cannot acquire any NEW locks

```
Phase 1 (Growing):   acquire locks
                     |
                     [Maximum Locking Point]
                     |
Phase 2 (Shrinking): release locks
```

2PL Compliance Example

T1a NOT 2PL ■	T1b 2PL ■	T2 2PL ■
lock_exclusive(a)	lock_exclusive(a)	lock_exclusive(a)
a = a - x	a = a - x	a = (1+r)a
unlock(a) <- releases	lock_exclusive(b)	lock_exclusive(b)
lock_exclusive(b) <- NEW	unlock(a)	b = (1+r)b
b = b + x	b = b + x	unlock(b)
unlock(b)	unlock(b)	unlock(a)
commit	commit	commit

```
T1a -> acquires new lock AFTER releasing -> violates 2PL ■
T1b -> acquires all locks BEFORE releasing -> 2PL ■
T2  -> acquires all locks BEFORE releasing -> 2PL ■
```

Schedule Sa vs Sb — Why it Matters

8. Deadlocks

- Deadlock = two transactions **waiting for each other** forever
- Neither can proceed → system stuck ■
- Called **circular wait**

```
T1 holds A, wants B --+
                        <-> circular wait ■
T2 holds B, wants A --+
```

Step by step:

```
T1: lock(A)      ■ gets A
T2: lock(B)      ■ gets B
T2: wants lock(A) ■ waiting... T1 has it
T1: wants lock(B) ■ waiting... T2 has it
-> DEADLOCK ■
```

Prevention

- Lock all objects in **same order** always

Both must lock: A first -> then B

```
T1: lock(A) --> lock(B) ■
T2: lock(A) --> lock(B) ■ (waits for T1 to release A)
-> No deadlock ■
```

Problems with prevention:

- Not always feasible ■
- Transactions often independent — don't know in advance what they'll need
- Reduces concurrency ■

Detection & Resolution

Detection — Wait-For Graph:

- Build graph of who is waiting for who
- Cycle in graph = deadlock

```

T1 --waits--> T2
^              |
+---waits-----+
      CYCLE = DEADLOCK ■

```

In distributed systems -> nodes share info -> detect cycle locally

Resolution — Abort a Victim:

- Abort one transaction (victim) to break the cycle
- Choose victim to **minimise rollback cost**

Strategy 1 -> abort YOUNGEST transaction
Strategy 2 -> abort transaction LOCAL to detection site
After victim aborts -> other transaction proceeds ■

9. Transactional Middleware

How it Started

- Originally → transactions managed **inside DBMS** (Persistence Tier)

Process 1 --+

Process 2 --+--> [DBMS manages transactions]

Process n --+

Problem

- Systems became **distributed**
- Transactions applied to non-DB services too (messaging etc.)
- DBMS alone not enough ■ → move to **Business Logic Tier (Middleware)**

OLD: Processes --> DBMS handles everything

NEW: Processes --> MIDDLEWARE manages --> DBMS stores

Two Key Components

Component	Role
Transaction Manager (TM)	Coordinates overall transaction across ALL resources
Resource Manager (RM)	Manages ONE local resource. Ensures ACID for that resource

```
      +--> [RM1 - Database]
[TM]  -----+--> [RM2 - Message Queue]
      +--> [RM3 - Legacy System]
```

10. JEE Architecture

- JEE = Java Enterprise Edition
- Platform for large scale distributed applications
- Organised into **3 Tiers**

```
Tier 1          Tier 2          Tier 3
-----          -----          -----
[Client] <--> [J2EE Server      ] <--> [Database]
              |- EJB Container      Server
              |  +- Enterprise Beans
              +- Web Container
                  |- JSP Files
                  +- Servlets
```

Tier 1 -> Client (browser, app)

Tier 2 -> Business Logic (J2EE Server)

Tier 3 -> Data Storage (Database)

11. EJB Container

- Enterprise Beans run inside **EJB Container**
- Container = runtime environment that **controls the beans**
- Provides all system-level services automatically

Service	What it does
---------	--------------

Transaction Management	Handles ACID properties
Security	Controls access
Remote Client Connectivity	Allows remote clients to connect
Life Cycle Management	Creates/destroys beans as needed
Database Connection Pool	Manages DB connections efficiently

12. Types of EJBs

Enterprise Beans

```
|
+--> Session Beans    (temporary, per client)
|
+--> Entity Beans     (permanent, shared data)
```

Session Beans

- Represents a **client** in J2EE server
- Uses **Facade Pattern** → front face for client
- One bean per client only — dies when client disconnects → **non-persistent**

Property	Detail
One client	Serves ONE client only
Transient	Dies when client disconnects
Non-persistent	Data NOT saved to DB
Lifetime	Same as client session

Entity Beans

- Represents a **business object** in persistent storage
- Examples: Customer, Product, Order
- Multiple clients can access — data saved **permanently**

Property	Detail
Persistent	Data saved permanently
Shared	Multiple clients can access
Storage	DB, object store, file, legacy app
Represents	Real business domain object

Two types of persistence:

BMP (Bean Managed Persistence)

- > Developer writes SQL manually
- > More control, more work


CMP (Container Managed Persistence)

- > Container handles SQL automatically
- > Less code, easier ■

EJB Application Example

```
Tier 1          Tier 2                      Tier 3
-----          -----                      -----
[Browser] <--> [Servlet - Web Container    ]
                [Shopping Session Bean    ] <--> [Database]
                [ +-- Product Entity Bean  ]
                [ +-- Customer Entity Bean ]
                [ +-- Order Entity Bean    ]
```

Browser --> Servlet --> Shopping Session Bean



FULL SUMMARY

Topic	Key Point
Transaction	Change between two stable states
Atomic Action	Indivisible — all or nothing
Composite Serialised	Sequence of atomic actions
Composite Concurrent	Multiple actions sharing one object
ACID	4 properties every transaction must have
Atomicity	All or nothing — rollback on failure
Consistency	Stable state -> stable state always
Isolation	Concurrent result = serial result
Durability	Commit = permanent, survives crashes
Concurrency Control	Ensures isolation via locking
Serialisability	Concurrent schedule = serial result
S1/S2/S3	Same transactions, different execution orders
Violations	Unrepeatable read, dirty read, lost write
Locking	Shared (read) or Exclusive (read+write)
2PL	Acquire ALL locks before releasing any
Sa vs Sb	Sa fails (non-2PL), Sb correct (2PL)
Deadlock	Circular wait — two transactions stuck
Prevention	Lock in same order always
Detection	Find cycle in wait-for graph
Resolution	Abort victim transaction
Transactional Middleware	TM coordinates, RMs manage local resources
JEE Architecture	3 tiers — Client, Business Logic, Database
EJB Container	Manages beans + provides all services
Session Bean	Per client — transient, non-persistent
Entity Bean	Shared data — persistent, permanent
BMP vs CMP	Manual SQL vs automatic SQL by container